

Evaluating Large Language Models for Personalised Credit Card Recommendation from User Spending Profiles

Hindvrat Singh Rao

A dissertation submitted in partial fulfilment of the requirements
for the degree of MSc Advanced Computer Science with Artificial Intelligence

Department of Computer and Information Sciences
University of Strathclyde

August 1, 2026

Declaration

This dissertation is submitted in part fulfilment of the requirements for the degree of MSc Advanced Computer Science with Artificial Intelligence of the University of Strathclyde.

I declare that this dissertation embodies the results of my own work and that it has been composed by myself. Following normal academic conventions, I have made due acknowledgement to the work of others.

This research did not involve human participants, real personal data, or any other activity requiring ethics approval, and therefore no application to the Departmental Ethics Committee was required (see Section 3.6 for the full ethical-considerations statement).

I give permission to the University of Strathclyde, Department of Computer and Information Sciences, to provide copies of the dissertation, at cost, to those who may in the future request a copy of the dissertation for private study or research. (please tick) Yes ☐ No ☐

I give permission to the University of Strathclyde, Department of Computer and Information Sciences, to place a copy of the dissertation in a publicly available archive. (please tick) Yes ☐ No ☐

I declare that the word count for this dissertation (excluding title page, declaration, abstract, acknowledgements, table of contents, list of illustrations, references and appendices) is **9,513** words at the time of this draft, to be confirmed at final submission.

I confirm that I wish this to be assessed as a **Type 5** Dissertation.

Signed: _____

Date: _____

Abstract

This dissertation evaluates whether a free-tier, open-weight Large Language Model (Meta’s Llama 3.1, 8B parameters, accessed via the Groq API) can generate personalised UK credit-card recommendations competitive with a conventional content-based recommender, and how prompting strategy affects recommendation quality, factual reliability, explanation quality, and consistency. Using a fixed catalogue of 20 real UK credit cards and 300 synthetic user profiles — spanning eight consumer archetypes with spending distributions calibrated against the ONS *Living Costs and Food Survey 2022–23* — three prompting strategies (zero-shot, structured, few-shot) were evaluated across 900 LLM inference calls, scored against a content-based cosine-similarity baseline using a four-criteria evaluation framework: recommendation-baseline overlap, deterministic factual correctness, LLM-as-judge explanation quality, and repeated-run consistency.

Prompting strategy was found to have a large, statistically significant effect on recommendation quality (Kruskal-Wallis $H = 60.70$, $p = 6.6 \times 10^{-14}$), but in a direction contrary to common expectations from the general prompt-engineering literature: zero-shot prompting significantly outperformed both structured and few-shot prompting on overlap with the baseline (mean overlap 0.687 vs. 0.517 and 0.539 respectively) and top-1 exact-match accuracy (0.307 vs. 0.077 and 0.023). However, structured prompting produced the most factually reliable responses (73.7% of responses entirely free of checkable factual errors, against 55.0% for zero-shot and 24.0% for few-shot), demonstrating that recommendation accuracy and factual trustworthiness are distinct properties that a single-metric evaluation would fail to distinguish. Strikingly, an independent LLM judge rated few-shot prompting’s explanations as the clearest and best-grounded of the three strategies (mean 2.907/3, against 2.730 for zero-shot and 2.650 for structured, $p < 0.0001$) despite few-shot being the least accurate and least factually reliable strategy on every other measure — demonstrating that perceived explanation quality is not a reliable proxy for correctness. Zero-shot was also numerically the most stable strategy under repeated querying (mean pairwise Jaccard similarity 0.826, against 0.808 for few-shot and 0.776 for structured), though this difference did not reach statistical significance on the sample tested. No single prompting strategy dominates across all evaluated criteria.

These findings suggest that prompting-strategy choice for LLM-based recommendation should be validated empirically against the specific task and model rather than assumed from general NLP benchmarks, and that recommendation-quality and factual-reliability metrics should be reported and optimised for separately in applications where the reliability of the model’s stated reasoning matters.

Acknowledgements

I would like to thank my dissertation supervisor, Dimitri Russinov, for their guidance and feedback throughout this project, particularly in shaping the baseline methodology and the scale of the experimental dataset.

Contents

Declaration	i
Abstract	ii
Acknowledgements	iii
1 Introduction	1
1.1 Background and Motivation	1
1.2 Research Problem	2
1.3 Research Questions	2
1.4 Research Design Summary	2
1.5 Dissertation Structure	3
2 Literature Review	4
2.1 Large Language Models as Recommender Systems	4
2.2 Prompt Engineering and Prompting Strategies	4
2.3 Content-Based Filtering as an Evaluation Baseline	5
2.4 Evaluating LLM Output: Quality, Consistency, and Factuality	6
2.5 Financial Product Recommendation	6
2.6 Critical Synthesis and Positioning of This Study	7
3 Methodology	8
3.1 Research Design	8
3.1.1 Quality Criteria for the Experimental Design	8
3.2 Dataset Construction	9
3.2.1 Credit Card Catalogue	9
3.2.2 Synthetic User Profiles	10
3.2.3 Datasets Considered and Rejected	10
3.3 Baseline Method	11
3.4 LLM Experimental Setup	12
3.4.1 Model and Infrastructure	12
3.4.2 Prompting Strategies	12

3.4.3	Evaluation Framework	13
3.4.4	Consistency Sampling Rationale	14
3.5	Data Analysis	14
3.6	Ethical Considerations	14
3.7	Limitations	15
4	Analysis and Findings	16
4.1	Overview	16
4.2	Recommendation Quality by Prompting Strategy	16
4.3	Per-Archetype Breakdown	18
4.4	Factual Correctness	18
4.5	Consistency	19
4.6	Explanation Quality	20
4.6.1	Individual-Level Validity Check	21
5	Discussion	22
5.1	Interpretation of Findings by Research Question	22
5.1.1	RQ1: Overlap with the Baseline	22
5.1.2	RQ2: Effect of Prompting Strategy	22
5.1.3	RQ3: Factual Reliability and Quality of Stated Reasoning	23
5.1.4	RQ4: Consistency	24
5.2	Relation to the Literature	24
5.3	Limitations	25
6	Recommendations and Conclusions	27
6.1	Summary of Contributions	27
6.2	Answering the Research Questions	28
6.3	Recommendations	28
6.4	Concluding Remarks	29
	References	30
A	Repository and Code	32
B	Prompt Templates	33
B.1	Zero-Shot Prompt	33
B.2	Few-Shot Prompt (worked examples only; rules and format identical to zero-shot above)	34
B.3	Structured Prompt	34
B.4	LLM-as-Judge Prompt	35

C	Credit Card Catalogue (Extract)	36
D	CSV Schema Reference	37

List of Figures

4.1	Distribution of overlap scores by prompting strategy	17
4.2	Combined recommendation-quality comparison across strategies	18

List of Tables

4.1	Recommendation-quality summary by prompting strategy (n=300 profiles per strategy)	16
4.2	Paired Wilcoxon signed-rank tests, per-profile mean overlap score	17
4.3	Mean overlap score and top-1 accuracy by archetype and strategy	18
4.4	Factual correctness by strategy	19
4.5	Consistency (mean pairwise Jaccard similarity across 3 runs) by strategy, n=24 profiles per strategy	20
4.6	Explanation-quality scores (LLM-as-judge, 1–3 scale) by strategy, n=300 per strategy	20
C.1	Credit card catalogue extract	36

Chapter 1

Introduction

1.1 Background and Motivation

The UK credit card market offers consumers a bewildering range of products, each differentiated by annual fees, reward structures, welcome offers, and category-specific cashback or points rates. Choosing a suitable card requires matching a card’s reward structure against an individual’s actual spending pattern — a task that is, in principle, well suited to automated recommendation. Traditional recommender systems address this kind of problem through collaborative or content-based filtering (Lops, de Gemmis and Semeraro, 2011; Pazzani and Billsus, 2007), and such systems are already used across banking and financial services to personalise product offers (Met, 2024).

The recent emergence of general-purpose Large Language Models (LLMs) offers a different route to the same goal. Rather than encoding a user’s preferences as a fixed feature vector and computing similarity against a catalogue, an LLM can be given a natural-language (or structured JSON) description of a user’s spending profile and a description of the available products, and asked directly to reason about which products are the best fit. This is attractive because it requires no bespoke model training, can incorporate qualitative context a numeric feature vector would discard, and can, in principle, explain its reasoning in natural language. LLMs have already been explored extensively as recommendation engines in the wider literature (Wu et al., 2024; Liu et al., 2024), and financial services specifically have begun exploring transformer-based personalisation.

However, using an LLM as a recommender introduces a new set of questions that a similarity calculation does not: how the request is phrased (prompting strategy) can materially change the output, the model may state facts about a product that are simply wrong (hallucination), and the same input may not reliably produce the same output on repeated queries. These are not concerns for a deterministic, formula-based baseline, but they are central to whether an LLM-based recommender could be trusted in a real product.

1.2 Research Problem

This study investigates whether a free-tier, open-weight LLM (Meta’s Llama 3.1, 8B-parameter variant, accessed via the Groq inference API) can generate personalised UK credit-card recommendations that are competitive with a conventional content-based baseline, and how the choice of prompting strategy affects the quality, consistency, and factual reliability of its output.

1.3 Research Questions

This dissertation addresses the following research questions:

1. To what extent do LLM-generated top-3 credit-card recommendations overlap with those produced by a content-based cosine-similarity baseline, across 300 synthetic UK consumer profiles spanning eight distinct spending archetypes?
2. Does the prompting strategy used (zero-shot, structured, or few-shot) produce a statistically significant difference in recommendation quality?
3. How does prompting strategy affect the model’s ability to produce accurate, checkable reasoning for its recommendations, and how factually reliable are the specific card details (fees, reward rates) the model cites?
4. How stable are the model’s recommendations when the same profile and strategy are queried repeatedly?

1.4 Research Design Summary

The study follows a controlled experimental design (Type 5, per the Department’s dissertation classification). A fixed catalogue of 20 real UK credit cards and 300 synthetic user profiles — generated across eight consumer archetypes with spending distributions calibrated against the ONS *Living Costs and Food Survey 2022–23* — form the basis for all experiments. A content-based cosine-similarity recommender, grounded in established recommender-systems literature (Lops, de Gemmis and Semeraro, 2011; Pazzani and Billsus, 2007), provides the reference point against which all LLM outputs are evaluated. The same 300 profiles are passed to the LLM under three prompting conditions — zero-shot, structured, and few-shot — producing 900 LLM inference calls in total. Outputs are scored on recommendation overlap against the baseline, top-1 exact-match accuracy, factual correctness of stated card details (checked deterministically against the card catalogue), explanation quality (LLM-as-judge), and stability across repeated runs. Full methodological detail is given in Chapter 3.

1.5 Dissertation Structure

Chapter 2 reviews the relevant literature across four areas: LLMs applied to recommendation tasks; prompt engineering strategies and their reported effect on task performance; content-based recommender systems as an evaluation baseline; and methods for evaluating LLM output quality, including LLM-as-judge approaches. Chapter 3 details the methodology: dataset construction, baseline design, LLM experimental setup, and evaluation framework. Chapter 4 presents the empirical results. Chapter 5 discusses the findings in relation to the research questions and existing literature, and considers the study’s limitations. Chapter 6 concludes with recommendations for future work.

Chapter 2

Literature Review

2.1 Large Language Models as Recommender Systems

Recommender systems have traditionally relied on collaborative filtering, content-based filtering, or hybrid approaches, each requiring a structured representation of users and items (Pazzani and Billsus, 2007). The arrival of general-purpose LLMs has prompted a substantial and rapidly growing body of work exploring whether these models can perform recommendation directly, without task-specific training. Liu et al. (2024) and Wu et al. (2024) both survey this emerging area, characterising LLM-based recommendation approaches along a spectrum from using an LLM purely as a feature encoder within a conventional pipeline, through to using the LLM as the recommender itself via prompting — the approach adopted in this study.

A recurring theme in this literature is that LLMs bring capabilities a conventional recommender lacks: they can incorporate free-text context, generalise to items or users outside the training distribution (zero-shot generalisation), and produce natural-language justifications for their choices (Liu et al., 2024). These properties are directly relevant to a domain like credit-card recommendation, where a user’s stated goals (“I want to build credit,” “I’m planning a large purchase”) are naturally expressed in language rather than as a fixed numeric feature. At the same time, the same surveys consistently flag three risks that motivate the evaluation framework adopted in this dissertation: (i) sensitivity to how the task is phrased in the prompt, (ii) hallucination of facts not present in the supplied context, and (iii) a lack of the determinism and auditability that a conventional, formula-based recommender guarantees by construction.

2.2 Prompt Engineering and Prompting Strategies

Since Brown et al. (2020) demonstrated that GPT-3 could perform new tasks from a handful of in-context examples without any parameter updates — “few-shot learning” — prompt design has become a first-class variable in how LLMs are applied to downstream tasks, rather than a fixed implementation detail. Schulhoff et al. (2024), in one of the most comprehensive tax-

onomies of prompting techniques to date, catalogue dozens of distinct prompting techniques for text-based tasks alone, underlining that “zero-shot” and “few-shot” are two points on a much larger design space that also includes explicit output-format constraints (referred to here as “structured” prompting).

Empirical comparisons of zero-shot against few-shot prompting report task-dependent results. Few-shot prompting has been shown to yield meaningful accuracy gains over zero-shot in several NLP classification tasks, with in-context examples providing the model with an implicit template for both reasoning and expected output format, and gains are typically achieved with as few as one to five examples. However, this advantage is not universal: for some domains, particularly complex reasoning and specialised tasks, well-specified zero-shot instructions have been found to be competitive with, or in some cases superior to, few-shot prompting, suggesting that the relative benefit of providing examples depends heavily on how much the task benefits from an explicit template versus additional worked examples. This inconsistency in the literature is itself a motivation for this dissertation’s comparative design: rather than assuming *a priori* that one strategy dominates, the three strategies are evaluated empirically against a common baseline and a common evaluation framework.

The “structured” prompting condition used in this dissertation — an explicit output schema (JSON, with fields for card ID, name, and justification) specified in the prompt — reflects a widely used mitigation for a separate, practical problem: unconstrained natural-language output is difficult to parse reliably. Requiring the model to populate a fixed schema is one of the more common patterns discussed in the prompting-technique literature (Schulhoff et al., 2024) precisely because it improves downstream parseability, which matters for an evaluation pipeline like the one used here that must automatically extract and score card recommendations from raw model output at scale (900 responses).

2.3 Content-Based Filtering as an Evaluation Baseline

Content-based recommendation constructs a profile of item attributes and a corresponding profile of user preferences over those same attributes, then recommends items whose profile most closely matches the user’s, typically via cosine similarity or a related distance metric (Pazzani and Billsus, 2007; Lops, de Gemmis and Semeraro, 2011). This family of methods is well established, fully deterministic, and — critically for this dissertation’s purposes — entirely independent of any language model, making it an appropriate conventional comparator for LLM-generated recommendations. This dissertation’s baseline (Section 3.3) follows this established approach directly: each of the 20 catalogue cards and each of the 300 synthetic profiles are encoded as numeric feature vectors over the same set of attributes (annual fee, category-specific reward rates, foreign-transaction fee), and cosine similarity between a profile vector and every card vector determines the top-3 reference recommendation for that profile.

2.4 Evaluating LLM Output: Quality, Consistency, and Factuality

Evaluating LLM-generated text is itself an active research area, distinct from evaluating a classification or ranking model’s accuracy, because LLM output is unstructured, open-ended, and often has no single “correct” answer to compare against. Zheng et al. (2023) introduced the “LLM-as-a-judge” paradigm — using a strong LLM to score the outputs of another model against a rubric — after showing that GPT-4 judgments agreed with crowdsourced human preferences over 80% of the time, comparable to human–human agreement, on their MT-Bench and Chatbot Arena benchmarks. They also documented specific failure modes of this approach relevant to any study using it, including position bias, verbosity bias, and self-enhancement bias. This dissertation’s evaluation design (Section 3.4.3) mitigates self-enhancement bias specifically by using a larger, different model (Llama 3.3 70B) to judge the outputs of the smaller model under evaluation (Llama 3.1 8B), rather than having a model judge its own output.

Two further evaluation properties matter specifically for a recommendation task rather than a general chat task. First, factual correctness — whether specific facts an LLM states about a recommended product are actually true — can, unlike open-ended explanation quality, be checked deterministically when a fixed, known catalogue of ground-truth facts exists, as is the case here (Section 3.4.3). Second, consistency — whether a model gives materially the same answer when queried repeatedly with the same input — is a property with no equivalent in a deterministic baseline, but is essential to any claim that an LLM-based recommender could be relied upon in practice; a system whose recommendation changes at random between identical requests is difficult to trust regardless of its average accuracy.

2.5 Financial Product Recommendation

Recommendation in financial services has typically relied on the same collaborative- and content-based methods used elsewhere, applied to transaction histories, account data, or stated preferences (Met, 2024). More recent work has explored deep learning and transformer-based approaches for financial product recommendation specifically, citing the need to capture more complex relationships between customer behaviour sequences and product attributes than simple similarity metrics allow. This dissertation sits at the intersection of this financial-recommendation literature and the LLM-recommendation literature reviewed in Section 2.1: it applies an LLM, rather than a purpose-trained deep learning model, to a financial product recommendation task, and evaluates it against the same kind of content-based baseline that has historically been the standard approach in this domain.

2.6 Critical Synthesis and Positioning of This Study

Read together, the four strands of literature reviewed above expose a structural gap rather than a purely empirical one. The LLM-recommendation surveys (Section 2.1) establish *that* LLMs can be used as recommenders and catalogue *how*, but are largely architecture- and technique-focused; they do not, in general, interrogate whether a given prompting strategy is actually reliable enough to trust in a specific, high-stakes consumer domain — the surveys treat prompting as one design choice among many rather than as a variable worth isolating and testing empirically against a domain-appropriate baseline. Conversely, the prompt-engineering literature (Section 2.2) isolates prompting strategy as its central variable, but the reported zero-shot/few-shot comparisons are drawn almost entirely from general NLP benchmarks — classification, sentiment, biomedical QA — not from recommendation tasks with a real product catalogue and a domain-specific correctness criterion. Neither strand, taken alone, answers the question this dissertation asks: for a bounded, factually-checkable recommendation task with real product data, does prompting strategy matter, and does it matter for the same reasons (accuracy) or for different ones (consistency, factual reliability) than in general NLP tasks?

This distinction matters methodologically. Most of the zero-shot/few-shot comparisons surveyed in Section 2.2 use accuracy or F1 against a single correct label as their only outcome measure. A recommendation task has no single correct answer in the same sense — the baseline in Section 2.3 provides a principled reference point, not a ground truth in the classification sense — which is precisely why this dissertation’s evaluation framework (Section 3.4.3) deliberately combines four distinct criteria (overlap with baseline, factual correctness, explanation quality, consistency) rather than a single accuracy figure. This is a direct methodological response to a limitation visible across the reviewed prompt-engineering literature: a model can be prompted into a *more accurate-looking* answer while becoming less consistent or less factually grounded, and a single-metric evaluation would not detect that trade-off. The results in Chapter 4 bear this out directly.

The financial-recommendation literature (Section 2.5) and the LLM-as-judge literature (Section 2.4) each fill a different part of the remaining gap: the former supplies the domain and the conventional baseline this dissertation compares against, while the latter supplies a defensible, literature-grounded method for scoring the one evaluation criterion (explanation quality) that cannot be checked deterministically. No source identified in this review combines all three elements — a real, factually-checkable product catalogue; a controlled, within-subjects comparison of prompting strategies; and a mixed quantitative/automated-qualitative evaluation framework covering accuracy, factuality, and stability together — which is the specific gap this dissertation addresses.

Chapter 3

Methodology

3.1 Research Design

This research follows a controlled experimental design, consistent with Type 5 (Experimental) dissertations under the Department’s classification. A fixed dataset of 300 synthetic user profiles and a curated catalogue of 20 real UK credit cards form the basis for all experiments. The same 300 profiles are evaluated under three prompting conditions, and outputs are assessed against a content-based recommender baseline. The evaluation combines quantitative scoring (recommendation overlap, top-1 accuracy, deterministic factual correctness) with an automated qualitative criterion (explanation quality, scored via LLM-as-judge — Section 3.4.3), producing a mixed-method analysis of LLM recommendation quality. The study is entirely computational: it does not involve human participants or real personal financial data (Section 3.6).

3.1.1 Quality Criteria for the Experimental Design

Following the Department’s guidance for Type 5 (Experimental) dissertations, this design is assessed against the four standard criteria for quantitative research design:

- **Replicability** — can the method be repeated? Every stage of the pipeline is fully specified and implemented in version-controlled, publicly available code (Appendix A): profile generation is deterministic (fixed random seed, Section 3.2.2), the baseline computation is a fully specified formula (Section 3.3), the exact prompt text for all three strategies is given in Appendix B, and model, temperature, and inference parameters are stated in Section 3.4.1. A third party with the same (free-tier) API access could reproduce this study end to end.
- **Reliability** — are the results repeatable, i.e. would a fresh sample from the same population show the same general pattern? This is assessed directly and quantitatively by the consistency criterion (Section 4.5): the same profile and strategy, queried three times

independently, shows moderate-to-high stability (mean Jaccard 0.78–0.83 across strategies), giving empirical evidence for reliability rather than assuming it.

- **Validity** — do the results have integrity, i.e. are the research instruments well-designed and appropriate? The baseline is grounded in an established, peer-reviewed method rather than a self-defined heuristic (Section 3.3); the LLM-judge uses a distinct, larger model than the one under evaluation specifically to reduce self-enhancement bias (Section 3.4.3). Both custom-built instruments were also directly validated rather than assumed correct: the factual-correctness checker (`src/factual_correctness.py`) was tested against a constructed set of test cases with known ground truth (correct claims, incorrect claims, and vague/unverifiable claims), correctly classifying 7 of 9 cases, with both misclassifications being conservative under-flags (a wrong claim scored as “uncertain” rather than “incorrect”) rather than false accusations against a correct claim — empirically confirming the checker’s own documented bias toward under- rather than over-counting errors. The LLM-judge was validated by checking whether its per-response scores actually correlate with independently-measured factual accuracy for the same response (Section 4.6.1); the result — a small, statistically significant *negative* correlation between judge score and factual correctness — is reported and discussed openly (Chapter 5) as a genuine limitation of the judge instrument uncovered by this validation, not concealed.
- **Generalisability** — can the results be applied beyond this exact context? Discussed fully in Section 5.3: the findings are specific to one model at one parameter scale (Llama 3.1, 8B), one card catalogue, and ONS-calibrated synthetic (not real) UK spending profiles, so they should be read as demonstrating a transferable *pattern* (that prompting strategy materially affects both accuracy and factual reliability, and that these two properties can trade off against each other) rather than as numeric findings that would reproduce exactly on a different model, catalogue, or population.

3.2 Dataset Construction

3.2.1 Credit Card Catalogue

A structured catalogue of 20 UK credit cards was assembled from publicly available issuer sources (see `docs/card-sourcing-notes.md` for provenance). Each entry records the card name, issuer, network, annual fee, reward type, reward rate by spending category (groceries, fuel, dining, travel, online shopping), welcome offer value, foreign transaction fee, and APR range. This catalogue defines the fixed product space from which all recommendations — both the baseline’s and the LLM’s — are drawn; no card outside this set can be a valid recommendation, which is itself checked as part of the evaluation pipeline (Section 3.4.3).

3.2.2 Synthetic User Profiles

All user profiles are synthetically generated; no real personal financial data is used at any stage (Section 3.6). An initial 80-profile dataset, using hand-specified archetype ranges, was expanded to 300 profiles following supervisor feedback that a larger, more rigorously grounded sample was needed for adequate statistical power in the planned significance testing.

Profiles are built around eight fixed archetypes representing distinct UK credit-card consumer segments: student (low spend), family (grocery-heavy), frequent traveller, balanced cashback user, large-purchase planner, online shopper, fuel commuter, and dining-heavy user. Rather than relying solely on hand-specified ranges, spending distributions across six categories (groceries, travel, fuel, dining, online shopping, other) are calibrated against the ONS *Living Costs and Food Survey 2022–23* (Office for National Statistics, 2023), so that each archetype’s spending pattern is anchored to an empirically observed UK household expenditure distribution rather than an arbitrary one (see `docs/data-justification.md` for the specific ONS figures used per category). This is consistent with the wider literature on financial-recommendation research, where individual-level real spending data is not publicly available for use due to UK GDPR and Financial Services and Markets Act 2000 restrictions, and simulated or anonymised profiles are the standard methodological choice.

Profiles are distributed 38/38/38/38/37/37/37/37 across the eight archetypes (summing to exactly 300), generated deterministically via a fixed random seed (`seed = 42`) for full reproducibility. Each profile contains: a profile identifier, archetype label, age group, monthly income, monthly spend across six categories, annual fee preference, primary goal, travel interest level, spending style, preferred reward type, a cost-sensitivity flag, and a planned large-purchase amount (non-zero only for the large-purchase-planner archetype).

3.2.3 Datasets Considered and Rejected

Before adopting the synthetic, ONS-calibrated approach described above, two categories of existing public dataset were examined as possible alternatives to constructing synthetic profiles, and both were found unsuitable for this study’s specific task. The first is the Office for National Statistics’ *Consumer Card Spending* series, published via Nomis and derived from aggregated, anonymised Visa card-transaction data (Office for National Statistics, 2024) — a separate ONS release from the *Family Spending* survey used for profile calibration in Section 3.2.2 above. This series is genuinely UK-sourced and real, but it is aggregated to postal-district level by merchant and cardholder location; it reports the flow of total spending between geographic regions, not individual-level records, and contains no per-person spending-category breakdown or credit-card product-preference information. It was therefore not usable for constructing individual user profiles.

The second category is individual-level transaction datasets published on data-sharing platforms such as Kaggle, for example the *Credit Card Transactions Dataset* (Choksi, 2024), which

contains approximately 1.85 million row-level transactions with per-transaction category, merchant, and amount fields. This is closer to the individual-level granularity this study’s task needs, but — like every such dataset examined — it contains no field indicating which credit-card product, or type of product, the cardholder actually holds or would prefer. A recommendation task needs exactly this link between a spending profile and a product-preference outcome, and no dataset examined provides it.

No publicly available dataset was found that combines individual-level spending by category with a genuine, linked credit-card product-preference label. This absence is unsurprising: assembling such a dataset from real individuals would itself constitute processing special-category financial data, restricted under UK GDPR and the Financial Services and Markets Act 2000 (Section 3.6). It is precisely this absence that motivates constructing synthetic profiles calibrated against real aggregate statistics (ONS, 2023) rather than either forcing an ill-fitting existing dataset onto the task or generating profiles with no empirical anchor at all.

3.3 Baseline Method

An initial baseline used four self-defined weighted heuristics (reward alignment, fee suitability, welcome-offer relevance, large-purchase suitability). Following supervisor feedback that this baseline was not grounded in established methodology, it was replaced with a content-based cosine-similarity approach, a well-established method in the recommender-systems literature (Lops, de Gemmis and Semeraro, 2011; Pazzani and Billsus, 2007; see Section 2.3).

Each credit card is represented as a numeric feature vector over eight attributes: annual fee (inverted, so a lower fee scores higher), base reward rate, category-specific reward rates (grocery, fuel, dining, travel, online shopping), and foreign transaction fee (inverted). Each user profile is represented as a preference vector over the same eight dimensions: an annual-fee-preference weighting derived from the profile’s stated fee tolerance, and category weightings derived from the profile’s monthly spend in each category (normalised to the profile’s own maximum category spend). Both card and profile vectors are min–max scaled to $[0, 1]$. For each profile, cosine similarity is computed between the profile vector and every card vector in the catalogue, and the three highest-scoring cards constitute the baseline recommendation, recorded in `data/ground_truth.csv`. This baseline is fully deterministic, transparent, reproducible, and independent of any language model, making it an appropriate conventional comparator for the LLM-generated recommendations evaluated in this study.

3.4 LLM Experimental Setup

3.4.1 Model and Infrastructure

The language model used to generate recommendations in this study is Meta’s Llama 3.1 (8B-parameter, instruction-tuned variant), accessed through the Groq inference API on its free tier. Groq was selected for its low-latency inference at no financial cost, appropriate for an academic project with resource constraints; Groq’s LPU (Language Processing Unit) infrastructure is purpose-built for LLM inference and offers sub-second time-to-first-token (Groq, 2024). The free tier used in this study is subject to a rate limit (requests per minute and tokens per minute) and a daily request cap, which materially affected the practical duration of the experimental runs (see Section 3.7). The model is used in standard inference mode with no fine-tuning or architectural modification; temperature was fixed at 0.2 across all calls to reduce (but not eliminate) run-to-run variance, with the residual variance itself the subject of the consistency evaluation (Section 3.4.4).

3.4.2 Prompting Strategies

Three prompt engineering strategies are evaluated across all 300 profiles:

- **Zero-shot prompting:** the model receives the user profile and the full card catalogue and is instructed to recommend the top three cards, with one short justification per card, without any worked examples.
- **Structured prompting:** the model receives the same input alongside an explicit JSON output schema specifying card ID, card name, a one-sentence reason, and any relevant warnings for each of the three ranked recommendations.
- **Few-shot prompting:** the model is provided with two complete worked examples — a profile alongside its correct recommendation and justification — before being presented with the target profile, to enable in-context learning (Brown et al., 2020; see Section 2.2).

All 300 profiles are evaluated under each of the three strategies, producing 900 LLM inference calls in total. Each result is stored with its profile identifier, strategy label, up to three extracted recommended card IDs, the full raw model response, and inference latency.

A methodological correction made during development: the zero-shot and few-shot prompt templates originally instructed the model to return bare card IDs only, with explicit instructions not to provide any explanation. This was identified as inconsistent with the evaluation framework’s “explanation quality” criterion (Section 3.4.3), which requires reasoning text to score. Both prompts were revised to require a one-sentence justification per recommended card, matching the reasoning already present in the structured condition, and the zero-shot and few-shot experimental runs were repeated under the corrected prompts. This correction was

completed successfully: all 900 responses (300 profiles $\times$ 3 strategies) were verified to contain zero duplicate profile–strategy pairs and, for zero-shot and few-shot, to be in the corrected reasoning-required format, before any scoring was performed.

3.4.3 Evaluation Framework

LLM outputs are assessed against four criteria:

1. **Recommendation quality** — the overlap between the LLM’s top-3 recommendation and the baseline’s top-3 recommendation for the same profile, reported as an overlap score (proportion of the three baseline cards also present in the LLM’s top three) and as top-1 exact-match accuracy.
2. **Factual correctness** — accuracy of specific, checkable card details cited in the response (card name, stated annual fee, stated reward-rate percentages), checked deterministically against `data/cards.csv`. Unlike the original methodology’s framing of this as a criterion requiring human judgement, a deterministic checker was implemented instead (`src/factual_correctness.py`): it extracts specific numeric/named claims from each response and compares them directly against ground truth, which is more reproducible and removes rater subjectivity for the subset of claims that are objectively checkable (see Section 3.7 for the scope and limits of this approach).
3. **Explanation quality** — clarity and apparent grounding of the model’s stated reasoning, scored via an LLM-as-judge approach (Zheng et al., 2023; Section 2.4) rather than the originally planned single-rater human scoring, given the scale of the dataset (900 responses) and the documented reliability of strong-model judgement against human preference (Zheng et al., 2023 report over 80% agreement, comparable to human–human agreement). The judge model is Llama 3.3 70B — deliberately larger than, and architecturally distinct in scale from, the Llama 3.1 8B model being evaluated, to reduce self-enhancement bias (Zheng et al., 2023). For each response, the judge scores each of the three recommendation lines and an overall explanation-quality score on a 1–3 ordinal scale (1 = poor: missing, generic, or irrelevant reasoning; 2 = adequate: understandable but vague or only partially grounded; 3 = good: clear and explicitly grounded in the user’s stated profile or the card’s actual attributes), returned as structured JSON for reliable parsing (`src/llm_judge.py`, `prompts/judge_explanation_quality.txt`). This substitution of model judgement for human judgement is treated explicitly as a methodological choice, not an oversight, and is discussed as a limitation in Section 3.7.
4. **Consistency** — stability of recommendations across repeated runs of the same profile and strategy, measured as the mean pairwise Jaccard similarity of the top-3 card sets across three runs, on a stratified subsample of profiles drawn evenly across all eight archetypes (`src/consistency_check.py`; see Section 3.4.4 for the sampling rationale).

Where the original methodology specified a 3-point ordinal rubric aggregated via Kruskal-Wallis testing across all four criteria, the two criteria with automated, continuous-valued scoring (recommendation quality, factual correctness) are analysed using their native continuous/proportion scales, with Kruskal-Wallis used specifically for the ordinal/count-based comparisons across the three prompt strategies, and paired Wilcoxon signed-rank tests used for pairwise strategy comparisons on a per-profile basis (matching the profile-level pairing inherent in the experimental design, i.e. every profile is evaluated under all three strategies).

3.4.4 Consistency Sampling Rationale

Repeating all 300 profiles across all three strategies would require 2,700 additional inference calls beyond the primary 900-call experiment. Given the free-tier rate and daily-request constraints observed during the primary data collection (Section 3.7), a stratified subsample is used instead: 3 profiles per archetype (24 profiles total, drawn with a fixed random seed for reproducibility) are each queried 3 times per strategy (216 additional calls in total), and stability is reported as the mean pairwise Jaccard overlap of the resulting top-3 sets. This ensures every archetype is represented in the consistency analysis without requiring a full-N repeat, a scope decision consistent with common practice in LLM stability studies where full repetition is cost-prohibitive.

3.5 Data Analysis

Scores are aggregated by prompt strategy. The Kruskal-Wallis H-test is used to test for an overall difference in recommendation-quality scores across the three strategies (appropriate for comparing more than two independent groups on non-normally-distributed ordinal/proportion data). Paired Wilcoxon signed-rank tests are additionally reported for each pairwise strategy comparison on a per-profile basis, since every profile is evaluated under all three conditions (a repeated-measures design), which the paired test is better powered to detect than the omnibus test alone. Per-archetype breakdowns are also reported, to check whether prompt-strategy effects are uniform across the eight consumer segments or archetype-dependent.

3.6 Ethical Considerations

This study does not involve human participants, real financial data, or personally identifiable information at any stage. All user profiles are entirely synthetic; the real UK consumer expenditure statistics (ONS, 2023) used for calibration inform only the statistical distribution of category-level spending values used to generate synthetic profiles, not any individual-level record. The study adheres to standard ethical guidelines applicable to computational and

data-science research and does not require formal ethical approval beyond departmental good-practice standards.

3.7 Limitations

Full discussion in Chapter 5; summarised here per the standard dissertation structure. The use of synthetic profiles means findings may not fully generalise to real consumer behaviour, even with ONS-calibrated distributions, since calibration does not capture irregular expenditure, life events, or behavioural nuances a real individual would exhibit. The study evaluates a single LLM (Llama 3.1, 8B parameters) for recommendation generation, which limits cross-model generalisability; a larger or differently-trained model may perform materially differently. Explanation-quality scoring via LLM-as-judge, while grounded in the literature (Zheng et al., 2023) and designed to mitigate self-enhancement bias through the use of a distinct, larger judge model, still substitutes model judgement for human judgement and inherits other documented biases of the approach (position and verbosity bias). The consistency evaluation uses a stratified subsample rather than the full 300 profiles, for the practical reasons given in Section 3.4.4. The deterministic factual-correctness checker (Section 3.4.3) can only verify claims that are specific and numeric or exactly named; vague or qualitative claims in a response are neither confirmed nor flagged, so the reported factual-error rate is a lower bound on true factual inaccuracy, not a complete audit. The free-tier API's rate and daily-request limits meant the primary 900-call run, and the correction rerun described in Section 4.2, each took multiple days of wall-clock time to complete, which is a practical constraint on this study's methodology rather than a property of the models or prompting strategies themselves.

Chapter 4

Analysis and Findings

4.1 Overview

This chapter reports results from the final, verified dataset: 900 LLM responses (300 profiles $\times$ 3 strategies), confirmed to contain zero duplicate profile–strategy pairs, with the zero-shot and few-shot strategies in the corrected reasoning-required format described in Section 3.4.2. Section 4.2 addresses Research Questions 1 and 2 (recommendation quality and the effect of prompting strategy); Section 4.3 reports the per-archetype breakdown; Section 4.4 addresses Research Question 3 (factual correctness); Section 4.6 reports on explanation quality and consistency (Research Questions 3 and 4).

4.2 Recommendation Quality by Prompting Strategy

Table 4.1 summarises the overlap score (proportion of the baseline’s top-3 cards also present in the LLM’s top three) and top-1 exact-match accuracy for each strategy, across all 300 profiles.

Table 4.1: Recommendation-quality summary by prompting strategy (n=300 profiles per strategy)

Strategy	Mean overlap	Median overlap	Std. dev.	Mean top-1 acc.	% zero-overlap
Zero-shot	0.687	0.667	0.258	0.307	6.0%
Structured	0.517	0.333	0.299	0.077	10.7%
Few-shot	0.539	0.667	0.324	0.023	18.7%

Contrary to expectations set by the wider few-shot-learning literature (Section 2.2), **zero-shot prompting produced the strongest recommendation-quality results on every measure**: the highest mean overlap with the baseline (0.687, against 0.517 for structured and 0.539 for few-shot), the highest top-1 exact-match accuracy by a wide margin (0.307, roughly four times the structured rate and thirteen times the few-shot rate), and the lowest proportion of profiles receiving zero overlapping recommendations (6.0%, against 18.7% for few-shot).

A Kruskal-Wallis test confirms this difference is statistically significant across the three strategies: $H = 60.70$, $p = 6.6 \times 10^{-14}$ (well below the $p < 0.05$ threshold). Pairwise paired Wilcoxon signed-rank tests, matched at the profile level (Table 4.2), show that zero-shot significantly outperforms both structured ($W = 6996.0$, $p < 0.0001$) and few-shot ($W = 10097.5$, $p < 0.0001$), while the difference between structured and few-shot narrowly fails to reach significance at the conventional threshold ($W = 19809.0$, $p = 0.0588$).

Table 4.2: Paired Wilcoxon signed-rank tests, per-profile mean overlap score

Comparison	W	p -value	Significant ($p < 0.05$)
Zero-shot vs. Structured	6996.0	< 0.0001	Yes
Zero-shot vs. Few-shot	10097.5	< 0.0001	Yes
Structured vs. Few-shot	19809.0	0.0588	No

Figure 4.1 shows the distribution of overlap scores by strategy, and Figure 4.2 shows the combined per-strategy comparison generated by the analysis pipeline.




Figure 4.1: Distribution of overlap scores by prompting strategy

This result is discussed in relation to the prompt-engineering literature in Section 5.1: it is a genuine, counter-to-expectation finding rather than a data artefact, and one plausible explanation — that the additional instruction text and worked examples in the structured and few-shot conditions crowd out the model’s attention on the profile’s specific attributes, at only 8 billion parameters — is explored further in Chapter 5.




Figure 4.2: Combined recommendation-quality comparison across strategies

4.3 Per-Archetype Breakdown

Table 4.3 reports mean overlap and top-1 accuracy separately for each of the eight consumer archetypes, to check whether the strategy effect found in Section 4.2 is uniform across consumer segments or archetype-dependent.

Table 4.3: Mean overlap score and top-1 accuracy by archetype and strategy

Archetype	Mean overlap			Mean top-1 accuracy		
	Zero-shot	Structured	Few-shot	Zero-shot	Structured	Few-shot
Balanced cashback	0.798	0.588	0.667	0.368	0.000	0.000
Dining-heavy	0.640	0.261	0.640	0.568	0.000	0.000
Family, grocery-heavy	0.728	0.579	0.684	0.000	0.000	0.000
Fuel commuter	0.991	0.982	1.000	0.081	0.081	0.000
Large-purchase planner	0.342	0.261	0.333	0.054	0.027	0.162
Online shopper	0.766	0.802	0.685	0.162	0.432	0.027
Student, low-spend	0.579	0.289	0.000	0.237	0.000	0.000
Frequent traveller	0.649	0.377	0.316	0.974	0.079	0.000

The zero-shot strategy’s advantage is not uniform: it is dramatic for the `traveller_frequent` archetype (97.4% top-1 accuracy against 7.9% for structured and 0% for few-shot) and for `dining_heavy` (56.8% against 0% for both other strategies), but the strategies converge closely for `fuel_commuter`, where all three achieve near-perfect overlap (0.98–1.00) — consistent with this archetype’s spending profile mapping onto a small, easily-identified subset of the catalogue regardless of prompting approach. The `student_low_spend` archetype is the clearest failure case for few-shot prompting specifically, with a mean overlap of exactly 0.000, meaning the few-shot condition never once recommended a baseline top-3 card for this archetype across all 38 profiles; this is discussed further in Section 5.1.

4.4 Factual Correctness

Table 4.4 summarises the deterministic factual-correctness check (Section 3.4.3) across all 900 responses.

Table 4.4: Factual correctness by strategy

Strategy	Responses w/ claims	Claims checked	Correct	Incorrect	Uncertain	% error-free
Zero-shot	170/300	1,800	864	771	165	55.0%
Structured	300/300	1,715	1,247	346	122	73.7%
Few-shot	232/300	3,150	1,021	2,117	12	24.0%

Structured prompting produces both the most consistently checkable responses (300/300 contain at least one checkable claim, against 170/300 for zero-shot and 232/300 for few-shot) and the highest accuracy rate on those claims (73.7% of structured responses are entirely error-free, against 55.0% for zero-shot). Few-shot prompting is a clear outlier in the opposite direction: despite generating the largest volume of checkable claims (3,150, roughly $1.8\times$ structured’s total), only 24.0% of few-shot responses are entirely free of factual errors, and 2,117 of the 3,150 claims checked (67.2%) were incorrect. This suggests the two worked examples in the few-shot prompt (Section 3.4.2) may lead the model to over-generate specific numeric claims about fees and reward rates by pattern-matching the examples’ level of detail, without those specific figures being reliably grounded in the actual target card’s data — a hypothesis explored further in Chapter 5.

Combined with Section 4.2, this produces a striking split rather than a single strategy dominating on every criterion: zero-shot leads on recommendation-baseline overlap, structured leads on factual reliability. Sections 4.5 and 4.6 add two further dimensions (consistency and explanation quality) to this picture, and — as shown below — the pattern becomes sharper still: no strategy leads on more than two of the four criteria, and few-shot in particular scores best on explanation quality while scoring worst on the two criteria that measure whether that explanation is actually accurate. This is precisely the kind of trade-off the four-criteria evaluation framework (Section 3.4.3) was designed to surface, and would not have been visible from recommendation-quality scoring alone.

4.5 Consistency

Table 4.5 reports the mean pairwise Jaccard similarity of the top-3 card set across three repeated runs, on the stratified 24-profile sample described in Section 3.4.4 (24 profiles $\times$ 3 strategies $\times$ 3 runs = 216 calls). A score of 1.0 means all three runs returned an identical top-3 set; 0.0 would mean no overlap at all between any pair of runs.

All three strategies show reasonably high stability on average (0.78–0.83 on the 0–1 scale), and 13 of 24 zero-shot profiles and 13 of 24 few-shot profiles were perfectly stable (identical top-3 across all three runs), against 9 of 24 for structured. Numerically, zero-shot is the most stable strategy and structured the least, which is notable alongside Section 4.4’s finding that structured is the *most* factually reliable strategy — a possible stability/reliability trade-off, discussed

Table 4.5: Consistency (mean pairwise Jaccard similarity across 3 runs) by strategy, n=24 profiles per strategy

Strategy	Mean	Std. dev.	Min	Max
Zero-shot	0.826	0.200	0.500	1.000
Few-shot	0.808	0.227	0.400	1.000
Structured	0.776	0.184	0.467	1.000

further in Chapter 5. However, a Kruskal-Wallis test across the three strategies on this metric is not statistically significant ($H = 0.78$, $p = 0.68$), so this difference should be read as a numerical trend on a modest stratified sample (24 profiles per strategy) rather than a confirmed effect — unlike the recommendation-quality and factual-correctness results in Sections 4.2 and 4.4, which are backed by the full 300-profile sample and are statistically significant.

The per-archetype pattern (not tabulated in full here; see `results/consistency_scores.csv`) is uneven: `fuel_commuter` and `traveller_frequent` are highly stable across all three strategies (0.78–1.00), while `online_shopper` and `dining_heavy` show more run-to-run variation, particularly for few-shot (0.42 and 0.67 respectively) — consistent with the pattern already seen in Section 4.3, where these strategies’ performance is generally less uniform across archetypes.

4.6 Explanation Quality

The explanation-quality criterion (LLM-as-judge, `src/llm_judge.py`, Section 3.4.3) scores each response 1–3 on the clarity and grounding of its stated reasoning, using Llama 3.3 70B as judge — a distinct, larger model than the one being evaluated, to reduce self-enhancement bias (Zheng et al., 2023). All 900 responses were judged. Table 4.6 summarises the results.

Table 4.6: Explanation-quality scores (LLM-as-judge, 1–3 scale) by strategy, n=300 per strategy

Strategy	Mean	Std. dev.	% scored 3/3	% scored ≤ 1
Few-shot	2.907	0.291	90.7%	0.0%
Zero-shot	2.730	0.459	73.7%	0.7%
Structured	2.650	0.478	65.0%	0.0%

The result here is, on its own, the opposite of the recommendation-quality and factual-correctness findings: **few-shot prompting produces the highest-rated explanations** (mean 2.907/3, with 90.7% of responses scoring the maximum), followed by zero-shot (2.730) and structured (2.650). A Kruskal-Wallis test confirms this is statistically significant ($H = 56.60$, $p = 5.1 \times 10^{-13}$), and all three pairwise Wilcoxon comparisons are also significant (zero-shot vs. structured $p = 0.016$; zero-shot vs. few-shot and structured vs. few-shot both $p < 0.0001$).

4.6.1 Individual-Level Validity Check

The strategy-level pattern above raises an obvious question: does the judge’s score for an individual response actually track that response’s factual accuracy, or is the aggregate pattern a coincidence of strategy-level style differences? This is checked directly by merging the explanation-quality scores with the factual-correctness results (Section 4.4) at the level of individual responses (n=702 responses with at least one checkable factual claim) and computing the correlation between judge score and factual accuracy for the same response. The result is a small but statistically significant *negative* correlation between judge score and being factually error-free (point-biserial $r = -0.173$, $p < 0.0001$; mean judge score 2.801 for responses containing at least one factual error, against 2.642 for entirely error-free responses), and a small positive correlation between judge score and the raw *number* of factual errors in a response (Spearman $\rho = 0.162$, $p < 0.0001$). In other words, the judge’s scores do not merely fail to track factual accuracy at the individual-response level — if anything, responses with more factual errors were rated very slightly *higher* on average. This individual-level result reinforces, rather than merely coincides with, the strategy-level finding above, and is discussed as direct evidence for a verbosity/confidence bias in Chapter 5.

Read alongside Sections 4.2 and 4.4, this produces the sharpest trade-off in the whole study: few-shot prompting writes the most highly-rated, clearest-*sounding* explanations, while simultaneously being the least accurate against the baseline (mean overlap 0.539, Section 4.2) and the least factually reliable (only 24.0% of responses entirely error-free, Section 4.4). In other words, on this task, a response reading as well-explained is not a reliable signal that it is either accurate or factually correct — the worked examples in the few-shot prompt appear to teach the model to write in a more confident, detailed, well-structured style, without a corresponding improvement (indeed with a measurable decline) in whether that style is backed by correct facts. This is discussed further in Chapter 5.

Chapter 5

Discussion

5.1 Interpretation of Findings by Research Question

5.1.1 RQ1: Overlap with the Baseline

Across all 300 profiles and three strategies, LLM-generated recommendations showed meaningful but partial agreement with the content-based cosine-similarity baseline (mean overlap ranging from 0.517 to 0.687 depending on strategy, against a maximum possible overlap of 1.0). This confirms that an 8-billion-parameter, free-tier LLM can approximate the output of a purpose-built content-based recommender to a substantial degree without any task-specific training, supporting the premise in Section 1.1 that LLMs are a viable alternative recommendation mechanism. It is not, however, evidence that the LLM and the baseline are interchangeable: an overlap of 0.687 at its highest (zero-shot) still means, on average, roughly one of the baseline’s three reference cards is absent from the LLM’s top three. Because the baseline itself is a fixed-formula approximation of “good fit” rather than ground truth in an absolute sense, a lower overlap score is not necessarily evidence of a worse recommendation — it may equally reflect the LLM incorporating profile information (a stated goal, a qualitative preference) that the baseline’s eight-dimensional feature vector cannot represent. This is an inherent limitation of using a deterministic baseline as the sole reference point, discussed further in Section 5.3.

5.1.2 RQ2: Effect of Prompting Strategy

The result that zero-shot prompting significantly outperformed both structured and few-shot prompting on recommendation-baseline overlap and top-1 accuracy (Section 4.2) runs counter to the general expectation, drawn from the wider prompt-engineering literature (Section 2.2), that providing worked examples (few-shot) or an explicit schema (structured) should improve task performance. Three plausible, non-exclusive explanations are considered here.

First, both the structured and few-shot prompts are substantially longer than the zero-shot prompt — the structured prompt adds a full JSON schema specification, and the few-shot

prompt adds two complete worked examples on top of that same instruction set. At 8 billion parameters, Llama 3.1 8B has comparatively limited capacity to maintain attention over long, multi-part instructions relative to larger frontier models; it is plausible that the extra instructional content in these two conditions displaces some of the model’s effective attention on the specific profile it is meant to be reasoning about, a phenomenon informally consistent with the “lost in the middle” effect reported for long-context prompting in other studies, though this study did not directly test that mechanism and it should be treated as a plausible interpretation rather than a demonstrated cause.

Second, the two worked examples in the few-shot prompt necessarily represent only two of the eight archetypes in this study’s dataset (a student-type profile and a frequent-traveller-type profile). The archetype-level breakdown in Table 4.2 shows the few-shot strategy performing at or near zero top-1 accuracy for several archetypes distant from the two worked examples (e.g. `family_grocery_heavy`, `balanced_cashback`), which is consistent with the model anchoring more heavily on the specific examples shown than on the general instruction, rather than generalising the underlying task.

Third, the factual-correctness results (Section 4.4) show few-shot prompting generating the largest volume of specific numeric claims while being the least accurate on those claims. This suggests the worked examples may be teaching the model a superficial pattern — “state a specific fee or reward-rate figure to sound authoritative” — rather than the underlying reasoning process, a documented risk of few-shot prompting when the examples are not sufficiently representative of the full input distribution.

These findings should not be read as a general claim that zero-shot prompting is superior to few-shot or structured prompting; they are specific to this model scale, this task, and this particular set of worked examples, and are precisely the kind of task-dependent result the mixed literature in Section 2.2 would predict.

5.1.3 RQ3: Factual Reliability and Quality of Stated Reasoning

The factual-correctness results (Section 4.4) show that requiring an explicit output schema (structured prompting) produced the most factually reliable responses (73.7% entirely error-free), even though structured prompting did not produce the best recommendation-quality scores. This is the clearest demonstration in this study’s results of why a single-metric evaluation would be inadequate for this task: a marker evaluating only recommendation overlap would conclude zero-shot is the best strategy outright, while a marker evaluating only factual reliability would reach the opposite conclusion for structured prompting. The four-criteria framework adopted in Section 3.4.3, motivated directly by the gap identified in the literature review (Section 2.6), is what makes this trade-off visible.

The explanation-quality results (Section 4.6) sharpen this picture further, and produce the single most important finding of this dissertation: **the strategy judged to write the best-**

quality explanations (few-shot, mean 2.907/3) is simultaneously the least factually reliable (24.0% error-free) and the least accurate against the baseline (mean overlap 0.539). A judge — whether an LLM, as here, or plausibly a human reader without access to ground truth — rating explanation quality from the text alone would rank few-shot’s output as the best of the three strategies. Only by separately, deterministically checking the specific factual claims against the real card catalogue does it become clear that this apparent quality is not matched by accuracy. This is a direct, empirical illustration of a risk that is easy to state abstractly but hard to demonstrate concretely: a fluent, well-structured, confident-sounding explanation is not evidence that an LLM’s recommendation is correct, and evaluating explanation quality in isolation — without an independent factual check — would have reached the wrong conclusion about which strategy to prefer.

5.1.4 RQ4: Consistency

The consistency results (Section 4.5) show zero-shot as numerically the most stable strategy (mean Jaccard 0.826) and structured as the least stable (0.776), with few-shot in between (0.808) — though, unlike the RQ2 and RQ3 results, this difference is not statistically significant on the 24-profile stratified sample used ($p = 0.68$), so it should be read as a trend rather than a confirmed effect. Taken together with RQ2 and RQ3, an interesting pattern emerges rather than one strategy simply being best: zero-shot is both the most accurate against the baseline *and* numerically the most consistent, which is internally coherent (a model reasoning more directly from the profile, without extra instructional content to weigh, might plausibly be both more accurate and more repeatable). Structured prompting, by contrast, is the most factually reliable but numerically the least consistent of the three — a response that gets the facts right more often is not necessarily the response that stays the same each time you ask. This is a genuine, if statistically tentative, three-way tension across the criteria rather than a single dominant strategy, and reinforces this dissertation’s central methodological argument (Section 2.6) that no single evaluation criterion tells the whole story for this kind of task.

5.2 Relation to the Literature

The finding that prompting strategy has a large, statistically significant effect on this task (Section 4.2) is consistent with the general premise of the prompt-engineering literature (Schulhoff et al., 2024) that prompt design is a first-class variable rather than an implementation detail. However, the specific direction of the effect — zero-shot outperforming few-shot — diverges from the more common finding in general NLP benchmarks that few-shot prompting improves performance (Section 2.2), reinforcing this dissertation’s positioning (Section 2.6) that recommendation tasks with a real, checkable product catalogue behave differently from the classification and QA benchmarks that dominate the existing zero-shot/few-shot comparison literature.

The factual-correctness results also extend the hallucination concerns raised generally in the LLM-recommendation survey literature (Liu et al., 2024; Wu et al., 2024) with a concrete, measured error rate specific to this domain, rather than a general acknowledgement that hallucination is possible.

The explanation-quality/factual-correctness divergence for few-shot prompting (Section 5.1) also connects directly to a specific, named risk in the LLM-as-judge literature: Zheng et al. (2023) document *verbosity bias*, where a judge (model or human) tends to rate longer, more detailed responses more favourably regardless of their actual correctness. The individual-level validity check in Section 4.6.1 provides direct evidence consistent with exactly this risk: judge scores correlate weakly but significantly *negatively* with individual responses being factually error-free ($r = -0.173$, $p < 0.0001$), meaning the judge model was not, in this study, a reliable proxy for factual accuracy even setting aside the strategy-level pattern — few-shot’s worked examples appear to have taught the model to write more detailed, specific-sounding justifications, and it is exactly this detail that the judge rates highly and that the factual checker subsequently shows to be frequently wrong. This is an argument for treating explanation-quality scores (from any judge, automated or human) as informative but insufficient on their own, and for pairing them with an independent, deterministic factual check wherever the underlying facts are checkable, as this study did.

5.3 Limitations

Synthetic data. All 300 user profiles are synthetic, calibrated against aggregate ONS expenditure statistics rather than drawn from real individual spending records (Section 3.2.2), because no such records are legally available for this kind of research (UK GDPR, Financial Services and Markets Act 2000). ONS calibration anchors the plausibility of the aggregate distributions used but cannot capture the irregular, bursty, or life-event-driven spending patterns a real consumer would exhibit, so findings about which prompting strategy performs best may not transfer perfectly to real users with messier spending histories.

Single model. All recommendations were generated by one model (Llama 3.1, 8B parameters). The explanation for the zero-shot finding proposed in Section 5.1 — that limited attention capacity at this model scale disadvantages longer prompts — is a testable hypothesis, not a confirmed result, because this study did not repeat the experiment with a larger model to check whether the effect persists, diminishes, or reverses at greater scale.

Baseline as reference point, not ground truth. As discussed in Section 5.1, the cosine-similarity baseline (Section 3.3) is itself a simplified, eight-dimensional approximation of “good fit,” not an independently validated gold standard. Overlap with the baseline measures agreement with one particular reasonable method, not objective recommendation quality in an absolute sense.

LLM-as-judge for explanation quality. Using a second LLM to score explanation quality

(Section 3.4.3) is grounded in the literature (Zheng et al., 2023) and specifically designed to reduce self-enhancement bias by using a distinct, larger judge model, but it still substitutes model judgement for the human judgement originally specified in the project’s methodology, and inherits other documented LLM-judge biases (position and verbosity bias) that a human rater would not share.

Deterministic factual-correctness scope. The factual-correctness checker (Section 3.4.3, `src/factual_correctness.py`) can only verify claims that are specific, numeric, or exactly named against the card catalogue; it cannot assess vague or qualitative claims (e.g. “this card offers good value”). The reported error rates (Section 4.4) are therefore a lower bound on the true rate of factual inaccuracy in the responses, not a complete audit of every claim made.

Consistency sampling scope. The consistency evaluation (Section 3.4.4) uses a stratified 24-profile subsample rather than all 300 profiles, for the practical rate-limit reasons given there; findings about stability generalise to the sampled archetypes but were not verified across every individual profile.

Infrastructure constraints on methodology. The free Groq API tier’s rate and daily-request limits (Section 3.4.1) meant that both the primary 900-call experiment and the subsequent zero-shot/few-shot correction rerun (Section 3.4.2) each took multiple days of wall-clock time to complete, constraining how many additional experimental conditions (e.g. additional models, additional prompting strategies, a full-N consistency check) could practically be included within the scope of this dissertation.

Chapter 6

Recommendations and Conclusions

6.1 Summary of Contributions

This dissertation set out to evaluate whether a free-tier, open-weight LLM (Llama 3.1, 8B parameters, via the Groq API) can generate personalised UK credit-card recommendations competitive with a conventional content-based baseline, and how prompting strategy affects the quality, factual reliability, explanation quality, and consistency of its output. Against a fixed catalogue of 20 real UK credit cards and 300 ONS-calibrated synthetic user profiles spanning eight consumer archetypes, three prompting strategies (zero-shot, structured, few-shot) were evaluated across 900 LLM inference calls, scored against a content-based cosine-similarity baseline and a four-criteria evaluation framework (recommendation quality, factual correctness, explanation quality, consistency).

The central empirical finding is that prompting strategy has a large, statistically significant effect on recommendation quality in this domain ($H = 60.70$, $p = 6.6 \times 10^{-14}$), but the effect runs in the opposite direction to what the general prompt-engineering literature would predict: zero-shot prompting outperformed both structured and few-shot prompting on recommendation-baseline overlap and top-1 accuracy. At the same time, structured prompting produced the most factually reliable responses (73.7% entirely error-free, against 55.0% for zero-shot and only 24.0% for few-shot), and few-shot prompting — despite generating the highest volume of specific factual claims — was the least reliable strategy on both recommendation quality and factual correctness. Most strikingly, few-shot prompting produced the *highest*-rated explanations by an independent LLM judge (mean 2.907/3, against 2.730 for zero-shot and 2.650 for structured, $p = 5.1 \times 10^{-13}$) despite being the least accurate and least factually reliable strategy on every other measure — demonstrating concretely that explanation quality, judged in isolation, is not a proxy for correctness. No single strategy dominates across every criterion, which is the core methodological argument this dissertation makes: a recommendation task evaluated on a single accuracy-like metric, or on explanation quality alone, would have reached a materially incomplete, and potentially misleading, conclusion about which prompting strategy is “best.”

6.2 Answering the Research Questions

- RQ1** LLM-generated recommendations showed substantial but partial overlap with the content-based baseline (mean overlap 0.517–0.687 depending on strategy), confirming an LLM can approximate a purpose-built recommender without task-specific training, while not fully replicating its output.
- RQ2** Prompting strategy produces a statistically significant difference in recommendation quality (Kruskal-Wallis $p < 0.0001$), with zero-shot significantly outperforming both other strategies (paired Wilcoxon $p < 0.0001$ in both comparisons) and structured vs. few-shot not reaching significance ($p = 0.0588$).
- RQ3** Factual reliability of stated reasoning varies substantially and independently of recommendation quality: structured prompting is most reliable (73.7% error-free), few-shot least (24.0% error-free). Explanation quality (LLM-judged, 1–3 scale) runs in the opposite direction to factual reliability: few-shot scores highest (2.907/3), structured lowest (2.650/3), both differences statistically significant. Recommendation accuracy, factual trustworthiness, and perceived explanation quality are therefore three separate properties of a response, not three views of the same thing, and must be evaluated separately.
- RQ4** Recommendation stability across repeated runs is numerically highest for zero-shot (mean pairwise Jaccard 0.826) and lowest for structured (0.776), with few-shot in between (0.808), though this difference is not statistically significant on the stratified sample tested (Kruskal-Wallis $p = 0.68$). Zero-shot’s numerical lead on both accuracy (RQ2) and stability (RQ4) is internally consistent; structured’s combination of high factual reliability (RQ3) but the lowest numerical stability (RQ4) is a genuine tension worth further investigation at greater sample size.

6.3 Recommendations

For practitioners considering an LLM-based recommender for a similarly bounded, factually-checkable product-recommendation task: this study’s results suggest that (i) prompting strategy should be treated as an empirical question to test against the specific task and model, not assumed from general-purpose benchmarks, since the direction of effect found here (zero-shot best) is the reverse of the common expectation; (ii) recommendation-quality, factual-reliability, and explanation-quality metrics should be evaluated and reported separately rather than assuming any one is a proxy for the others, since a strategy that looks best on one can be worst on another, as shown starkly by few-shot prompting scoring best on explanation quality while scoring worst on both accuracy and factual reliability in this study; (iii) if factual reliability of the stated reasoning matters for a given application (e.g. a live consumer-facing recommender,

where a factually wrong claim about a fee or reward rate is a real-world liability), structured output prompting is worth the trade-off in raw recommendation-overlap performance found here; and (iv) explanation quality alone — whether judged by an LLM or a human reader without access to ground truth — should not be used as a stand-in for factual accuracy when a deterministic check against real product data is possible, since this study found the two to move in opposite directions.

For future work building directly on this dissertation: the model-scale hypothesis proposed in Section 5.1.2 (that limited attention capacity at 8B parameters specifically disadvantages longer structured/few-shot prompts) could be tested directly by repeating this exact experimental design with a larger model (e.g. a 70B-parameter model, also available on Groq’s free tier) to see whether the zero-shot advantage persists, narrows, or reverses at scale. The consistency evaluation could be extended from the stratified 24-profile subsample used here to the full 300-profile set, infrastructure permitting. Finally, the deterministic factual-correctness checker built for this study (`src/factual_correctness.py`) is itself a reusable contribution: it could be extended to a wider card catalogue or a different financial-product domain (e.g. savings accounts, insurance) with comparatively little modification, since its core method — extracting named/numeric claims via pattern matching and checking them against a structured ground-truth table — is domain-agnostic.

6.4 Concluding Remarks

This dissertation demonstrates that a free, open-weight LLM can perform a real, factually-grounded personalised recommendation task at a level worth taking seriously as an alternative to conventional content-based filtering, while also demonstrating, empirically and specifically for this task, why prompting-strategy choice cannot be treated as a minor implementation detail: it is the single largest source of variation in output quality found in this study, and it affects different quality dimensions — accuracy against a baseline, factual reliability, and perceived explanation quality — in different, sometimes directly opposing, directions. The clearest illustration of this is also the dissertation’s central warning for anyone building or evaluating LLM-based recommenders: the strategy whose explanations read best to an independent judge was, on the same data, the strategy least likely to be right.

References

Brown, T. et al. (2020) ‘Language Models are Few-Shot Learners’, *Advances in Neural Information Processing Systems*, 33. Available at: <https://arxiv.org/abs/2005.14165>

Choksi, P. (2024) *Credit Card Transactions Dataset*. Kaggle. Available at: <https://www.kaggle.com/datasets/priyamchoksi/credit-card-transactions-dataset>

Groq (2024) ‘Groq Launches Meta’s Llama 3 Instruct AI Models on LPU Inference Engine’. Available at: <https://groq.com/blog/12-hours-later-groq-is-running-llama-3>

Liu, Q. et al. (2024) ‘Large Language Model Enhanced Recommender Systems: A Survey’. Available at: <https://arxiv.org/abs/2412.13432>

Lops, P., de Gemmis, M. and Semeraro, G. (2011) ‘Content-Based Recommender Systems: State of the Art and Trends’, in *Recommender Systems Handbook*. Boston, MA: Springer, pp. 73–105.

Met, İ. (2024) ‘Product Recommendation System With Machine Learning Algorithms for SME Banking’, *International Journal of Intelligent Systems*, 2024. Available at: <https://onlinelibrary.wiley.com/doi/10.1155/2024/5585575>

Office for National Statistics (2023) *Family Spending in the UK: April 2022 to March 2023*. Available at: <https://www.ons.gov.uk/peoplepopulationandcommunity/personalandhouseholdfinances/expenditure/bulletins/familyspendingintheapril2022tomarch2023>

Office for National Statistics (2024) *Consumer Card Spending, Flow of Spending Across the UK: 2019 to 2023*. Available at: <https://www.ons.gov.uk/economy/economicoutputandoutput/articles/consumercardspendingflowofspendingacrosstheuk/2019to2023>

Pazzani, M.J. and Billsus, D. (2007) ‘Content-Based Recommendation Systems’, in *The Adaptive Web*. Berlin: Springer, LNCS 4321, pp. 325–341.

Schulhoff, S. et al. (2024) ‘The Prompt Report: A Systematic Survey of Prompting Techniques’. Available at: <https://arxiv.org/abs/2406.06608>

Wu, L. et al. (2024) ‘A Survey on Large Language Models for Recommendation’, *World Wide Web*, 27(60). Available at: <https://dl.acm.org/doi/10.1007/s11280-024-01291-2>

Zheng, L. et al. (2023) ‘Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena’, *Advances in Neural Information Processing Systems*, 36. Available at: <https://arxiv.org/abs/2306.05685>

Appendix A

Repository and Code

The full source code, data, prompts, and generated results for this dissertation are available at:

<https://github.com/hindvratsingh-ux/dissertation-creditcard-llm>

The repository includes: `src/build_profiles.py` (synthetic profile generation), `src/baseline.py` (cosine-similarity baseline), `src/llm_eval.py` (the 900-call LLM experiment), `src/score_outputs.py` and `src/analyze_results.py` (recommendation-quality scoring and statistical testing), `src/factual_correctness.py` (deterministic factual-claim checking), `src/llm_judge.py` (LLM-as-judge explanation-quality scoring), `src/consistency_check.py` (repeated-run stability scoring), and `src/generate_plots.py` (figures). `docs/status-and-gap-analysis.md` in the repository documents the full development history and methodological corrections made over the course of the project.

Appendix B

Prompt Templates

B.1 Zero-Shot Prompt

You are a UK credit card recommendation expert.

Below is a user's spending profile as a JSON object, followed by the full credit card catalogue.

Your task: recommend exactly the top 3 credit cards that best match this user's spending habits, goals, annual fee preference, and reward type preference.

Rules:

- Only recommend cards that exist in the catalogue below. Do NOT invent card IDs.
- Return exactly 3 lines, one per recommended card, ranked best first.
- Each line MUST use this exact format:
CC001 - <one short sentence explaining why this card fits this user, citing the specific attribute (fee, reward rate, or offer) that justifies it>
- Do not include any text before the first line or after the third line. No headings, no extra commentary.

User Profile:

{profile}

Credit Card Catalogue:

```
{cards_csv}
```

B.2 Few-Shot Prompt (worked examples only; rules and format identical to zero-shot above)

--- EXAMPLE 1 ---

User Profile:

```
{"profile_id": "EX001", "profile_type": "student_low_spend",  
  "preferred_reward_type": "None", "annual_fee_preference": "low",  
  "goal": "build_credit", ...}
```

Recommendation:

CC004 - No annual fee and no rewards complexity, designed specifically for building credit history.

CC005 - No annual fee, simple entry-level card suited to a low, irregular spender.

CC011 - Student-focused card with no fee, matching the low-income/no-vehicle profile.

--- EXAMPLE 2 ---

User Profile:

```
{"profile_id": "EX002", "profile_type": "frequent_traveller",  
  "preferred_reward_type": "Points", "annual_fee_preference": "high",  
  "goal": "travel_rewards", ...}
```

Recommendation:

CC008 - Highest travel reward rate (1 point per £1 on travel) and points transfer to airlines, matching the high travel spend and points preference.

CC007 - Avios earned on BA spend directly rewards frequent flying, aligned with travel_rewards goal.

CC013 - Travel-focused points card with free travel insurance, complementing heavy travel spend.

B.3 Structured Prompt

The structured condition uses the same rules as zero-shot above, but specifies a JSON output schema with fields for card ID, card name, a one-sentence reason, and any relevant warnings, for each of the three ranked recommendations. See `prompts/structured.txt` in the repository for the full text.

B.4 LLM-as-Judge Prompt

See `prompts/judge_explanation_quality.txt` in the repository for the full explanation-quality scoring rubric given to the judge model (Section 3.4.3).

Appendix C

Credit Card Catalogue (Extract)

Table C.1 shows the first five entries of the 20-card catalogue used throughout this study (full catalogue in `data/cards.csv`).

Table C.1: Credit card catalogue extract

ID	Card Name	Issuer	Annual Fee (£)	Reward Type
CC001	American Express Platinum Cashback Everyday	American Express	0	Cashback
CC002	American Express Platinum Cashback	American Express	25	Cashback

Appendix D

CSV Schema Reference

Full column-level documentation for `data/cards.csv`, `data/profiles.csv`, `data/ground_truth` and `data/llm_results.csv` is maintained in `docs/schema-notes.md` in the repository, and is not reproduced in full here for brevity.